

Complete Detailed Summary: Attacking ECDSA with Nonce Leakage by Lattice Sieving

□ Paper Overview (Quick Context)

What is this paper about? This paper presents a new method to break the ECDSA digital signature system when small pieces of secret information (called “nonces”) leak out during the signing process. The researchers developed an improved attack that works even when only 1 bit of information leaks per signature.

Why is this important? ECDSA is widely used to secure digital communications (like Bitcoin transactions, secure messaging, etc.). If attackers can break it with minimal information leakage, it poses serious security risks.

SECTION-BY-SECTION DETAILED EXPLANATION

1. TITLE AND AUTHORS

Title: “Attacking ECDSA with Nonce Leakage by Lattice Sieving: Bridging the Gap with Fourier Analysis-based Attacks”

Simple Explanation: - **ECDSA** = Elliptic Curve Digital Signature Algorithm (a way to create digital signatures) - **Nonce Leakage** = When secret random numbers used in signing accidentally get revealed (even partially) - **Lattice Sieving** = A mathematical technique to find patterns in high-dimensional spaces - **Fourier Analysis-based Attacks** = Another mathematical approach using frequency analysis

What the title means: The researchers are presenting a new attack method that combines ideas from two different mathematical approaches to break ECDSA more effectively when small amounts of secret information leak.

2. ABSTRACT (Summary of the Entire Paper)

Let me break down the abstract paragraph by paragraph:

Paragraph 1: The Problem Context

Original Text Simplified: “The Hidden Number Problem (HNP) is used to attack cryptographic systems like ECDSA when some information about secret numbers leaks.”

Purpose: - Introduces the mathematical problem at the heart of these attacks - HNP = trying to find a hidden secret number when you have partial clues

Two Main Approaches: 1. **Lattice-based attacks** - Work well when you know longer pieces of the secret 2. **Fourier analysis-based attacks** - Work better when you know very little (like 1 bit)

The Challenge: - Lattice attacks struggle when only tiny amounts leak (like 1 bit per signature) - Lattice attacks also fail when the leaked information contains errors

Paragraph 2: The Solution (What This Paper Contributions)

Key Innovation: The researchers introduce a “parameter x ” that modifies existing lattice structures to make them work better with small leakages.

Technical Achievement: - Reduces the “lattice dimension” (makes the problem simpler) - The reduction is approximately $(\log l x)/l$, where l = number of leaked bits

New Methods Introduced: 1. **Interval reduction algorithm** - Narrows down where the secret could be 2. **Several predicates** - Rules to filter out wrong answers 3. **Pre-screening technique** - Eliminates bad candidates early

Extension: They also made their method work even when the leaked information has errors (noise).

Paragraph 3: The Results

Records Achieved: - Successfully attacked ECDSA with **1-bit leakage** on a **160-bit curve** - Even worked with **less than 1-bit leakage** (meaning very noisy/partial information) - Previous best lattice attack for 1-bit leakage only worked on a **112-bit curve** (much weaker)

Why This Matters: 160-bit curves are used in real-world systems, so this shows a practical threat.

3. INTRODUCTION

Paragraph 1: Background on the Hidden Number Problem

Purpose: Establish the historical context and importance of HNP

Simple Explanation: - HNP was created by two researchers (Boneh and Venkatesan) to study the security of Diffie-Hellman (a key exchange method) - Later, other researchers (Nguyen and Shparlinski) applied it to ECDSA - The key insight: If an attacker learns small pieces of the random numbers (nonces) used in ECDSA signatures, they can recover the entire secret key

The Attack Scenario: 1. ECDSA generates signatures using secret random numbers (nonces) 2. If these nonces partially leak (through side-channels, timing attacks, etc.) 3. An attacker can collect many signatures with leaked nonce information 4. By solving the HNP, they can recover the secret key

Paragraph 2: Fourier Analysis-Based Attacks

Purpose: Explain the first type of attack approach

What It Does: Uses mathematical techniques from signal processing (Fourier analysis) to find patterns in the leaked information.

Methodology: - Analyzes frequencies and patterns in the leaked bits - Works by transforming the problem into the frequency domain

Advantages: - Works well with very small leakages (even less than 1 bit) - Can handle errors in the leaked information

Limitations: - Requires many signatures (large sample size) - Very high computational cost (takes a lot of processing power) - Can take days or weeks to complete

State-of-the-Art: - Aranha et al. (2021) broke 192-bit ECDSA with less than 1-bit leakage

- But required massive computational resources

Paragraph 3: Lattice-Based Attacks

Purpose: Explain the second type of attack approach

What It Does: Converts the HNP into a geometric problem in high-dimensional space.

Methodology - The Transformation Chain: 1. HNP → BDD (Bounded Distance Decoding) → CVP (Closest Vector Problem) → uSVP (Unique Shortest Vector Problem)

Simple Analogy: Imagine you have a 3D grid of points (a lattice). The secret key corresponds to a special point in this grid. The attack tries to find the shortest vector (arrow) in this grid, which reveals the secret.

How It Works: - Uses “Kannan’s embedding” - a mathematical trick to transform the problem - Success depends on whether the target vector (the secret) is “short enough” - Uses algorithms like BKZ (Block Korkine-Zolotarev) and lattice sieving

Advantages: - Much faster than Fourier attacks when enough bits leak - Requires fewer signatures

Limitations: - Struggles when only 1 bit leaks per signature - Performance degrades with errors in leaked data - Previous best result: 112-bit curve with 1-bit leakage

The Belief: Many researchers thought lattice attacks were fundamentally limited for 1-bit leakage scenarios.

Paragraph 4: Recent Improvements

Purpose: Discuss the state-of-the-art before this paper

Key Development: - Xu et al. (2023) introduced “lattice sieving with predicate” - Used G6K (a lattice sieving library) with GPU acceleration - Improved results but still limited to 112-bit curves for 1-bit leakage

The Limitation: Even with these improvements, there was still a significant gap between lattice attacks and Fourier attacks for small leakages.

Paragraph 5: The Research Question

Purpose: State the open problem this paper addresses

The Question: Can we bridge the gap between lattice-based and Fourier analysis-based attacks?

What This Means: - Can lattice attacks be made to work as well as Fourier attacks for small leakages? - Can we get the speed benefits of lattice attacks while handling 1-bit leakage?

Paragraph 6-7: This Paper’s Contributions

Purpose: Outline what the paper achieves

Main Contributions:

1. **Algorithmic Tradeoff:**
 - Introduces parameter x to control lattice dimension
 - Smaller dimension = easier problem but requires more careful analysis
 - This is the key innovation
2. **Interval Reduction Algorithm:**
 - Narrows down the search space for the secret
 - Works by analyzing the constraints from leaked bits

3. **New Predicates:**
 - Filtering rules to eliminate wrong candidates
 - Based on geometric properties of lattice vectors
4. **Pre-screening Technique:**
 - Checks candidates early before expensive computations
 - Saves time by avoiding unnecessary work
5. **Error Handling:**
 - Extended methods to work with noisy/erroneous leakage
 - Important for real-world scenarios where leakage isn't perfect

Experimental Results:

Record 1: 1-bit Leakage on 160-bit Curve - First successful lattice attack at this level - Previous best: 112-bit curve - Shows 43% improvement in curve size

Record 2: Less than 1-bit Leakage - Successfully attacked with 0.75 bits of leakage per signature - Demonstrates robustness of the method

Record 3: Error Tolerance - Worked even when 5% of the leaked bits were wrong - Shows practical applicability

Comparison with Previous Work: - Outperforms Xu et al.'s method - Requires fewer signatures - Faster computation time - Handles larger curve sizes

4. PRELIMINARIES (Background Knowledge)

This section provides the mathematical foundation needed to understand the attack.

Section 2.1: Notations

Purpose: Define the mathematical symbols used throughout the paper

Key Notations Explained:

- \mathbb{Z} = Integers (whole numbers: ..., -2, -1, 0, 1, 2, ...)
- \mathbb{R} = Real numbers (all numbers including decimals)
- \mathbb{Z}_p = Integers modulo p (numbers from 0 to p-1 with wraparound)
- $\|v\|$ = Length (norm) of vector v
- $u \cdot v$ = Dot product of vectors u and v
- $[a, b]$ = Interval from a to b (all numbers between a and b)

Why This Matters: These notations are used to describe the mathematical operations in the attack precisely.

Section 2.2: Lattices

Purpose: Explain what lattices are and their basic properties

What is a Lattice?

Simple Explanation: A lattice is like a regular grid of points in space. Imagine a 2D grid of dots on graph paper, but extended to many dimensions.

Formal Definition: Given n independent vectors (b_1, b_2, \dots, b_n), a lattice L is all possible combinations: $L = \{z_1 b_1 + z_2 b_2 + \dots + z_n b_n : z_1, z_2, \dots, z_n \text{ are integers}\}$

Example in 2D: If $b_1 = (1, 0)$ and $b_2 = (0, 1)$, the lattice is all points with integer coordinates.

Key Concepts:

1. **Basis:**
 - The set of vectors $\{b_1, b_2, \dots, b_n\}$ that generate the lattice

- Like the “building blocks” of the lattice
 - Many different bases can generate the same lattice
2. **Lattice Dimension:**
 - The number of basis vectors (n)
 - Higher dimension = more complex lattice = harder problems
 3. **Determinant ($\det(L)$):**
 - A measure of the “density” of the lattice
 - Smaller determinant = points are closer together

Why Lattices Matter for Cryptography: Many hard problems in cryptography can be translated into problems about finding short vectors in lattices.

Section 2.3: Hard Problems on Lattices

Purpose: Introduce the computational problems that make lattices useful for cryptography

Problem 1: Shortest Vector Problem (SVP)

Simple Explanation: Find the shortest non-zero vector in the lattice.

Analogy: Imagine a 3D grid of points. Starting from the origin $(0,0,0)$, which grid point is closest to you (but not the origin itself)?

Why It's Hard: In high dimensions (100+), this becomes extremely difficult to solve exactly.

Approximate SVP: Instead of finding THE shortest vector, find a vector that's “close enough” (within a factor γ of the shortest).

Problem 2: Closest Vector Problem (CVP)

Simple Explanation: Given a target point t (not in the lattice), find the lattice point closest to t .

Analogy: You're standing somewhere between grid points. Which grid point is nearest to you?

Why It's Important: The HNP attack converts to solving CVP instances.

Problem 3: Bounded Distance Decoding (BDD)

Simple Explanation: A special case of CVP where you know the target is close to some lattice point (within a certain distance).

The Guarantee: If you know the target is within distance d of some lattice point, and d is small enough, you can find that lattice point efficiently.

Connection to HNP: When we have partial nonce leakage, we can construct a BDD instance where solving it reveals the secret key.

Section 2.4: Lattice Reduction Algorithms

Purpose: Explain the tools used to solve lattice problems

What is Lattice Reduction?

Simple Explanation: Given a “bad” basis (with long, nearly parallel vectors), find a “good” basis (with short, nearly orthogonal vectors) for the same lattice.

Analogy: Imagine describing a grid using two nearly identical arrows pointing in almost the same direction (bad basis). Lattice reduction finds two short, perpendicular arrows that describe the same grid (good basis).

Algorithm 1: LLL (Lenstra-Lenstra-Lovász)

What It Does: Finds a reasonably good basis quickly.

Properties: - Runs in polynomial time (fast) - Finds vectors within a factor of $2^{(n/2)}$ of the shortest vector - Good for small dimensions

When to Use: Quick preprocessing or when exact solutions aren't needed.

Algorithm 2: BKZ (Block Korkine-Zolotarev)

What It Does: Finds better bases than LLL by working on "blocks" of the basis.

How It Works: 1. Divide the basis into overlapping blocks 2. Find the shortest vector in each block (using SVP solvers) 3. Update the basis 4. Repeat until no improvement

Parameters: - Block size β : Larger β = better reduction but slower - Typical values: $\beta = 20$ to 80

When to Use: When you need better reduction than LLL and can afford more computation time.

Algorithm 3: Lattice Sieving

What It Does: Finds very short vectors by systematically searching through vector combinations.

How It Works (Simplified): 1. Start with a large list of lattice vectors 2. Combine pairs of vectors to create shorter vectors 3. Keep only the shortest vectors 4. Repeat until you can't find shorter vectors

Modern Implementations: - G6K (General Sieve Kernel): A highly optimized sieving library - Uses advanced techniques like "pump and jump" - Can use GPUs for acceleration

Advantages: - Finds shorter vectors than BKZ - Better for solving hard instances

Disadvantages: - Much slower than BKZ - Requires more memory - Exponential time complexity: $O(2^{(0.292n)})$

When to Use: When BKZ isn't strong enough and you need to find very short vectors.

Section 2.5: ECDSA (Elliptic Curve Digital Signature Algorithm)

Purpose: Explain how ECDSA works and where the vulnerability lies

What is ECDSA?

Simple Explanation: A method to create digital signatures using elliptic curves. It proves that you have a secret key without revealing the key itself.

Components:

1. **Domain Parameters:**
 - **p**: A large prime number defining the field
 - **E**: An elliptic curve equation
 - **G**: A base point on the curve
 - **n**: The order of G (how many points in the subgroup)
2. **Keys:**
 - **Secret key (α)**: A random number between 1 and $n-1$ (kept secret)
 - **Public key (Q)**: $Q = \alpha G$ (shared publicly)

Signing Process (Step-by-Step):

Input: Message m , secret key α

Step 1: Hash the message - Compute $e = H(m)$ where H is a hash function (like SHA-

256) - Convert e to an integer

Step 2: Generate random nonce - Pick a random number k (the nonce) from 1 to n-1 -
CRITICAL: This must be truly random and secret!

Step 3: Compute r - Calculate point $(x, y) = kG$ - Set $r = x \bmod n$ - If $r = 0$, go back to Step 2

Step 4: Compute s - Calculate $s = k^{-1}(e + \alpha r) \bmod n$ - If $s = 0$, go back to Step 2

Step 5: Output signature - The signature is (r, s)

Verification Process:

Input: Message m, signature (r, s) , public key Q

Step 1: Check validity - Verify that r and s are in $[1, n-1]$

Step 2: Compute hash - $e = H(m)$

Step 3: Compute values - $w = s^{-1} \bmod n$ - $u_1 = ew \bmod n$ - $u_2 = rw \bmod n$

Step 4: Compute point - $(x, y) = u_1 G + u_2 Q$

Step 5: Verify - Accept if $r = x \bmod n$, reject otherwise

Why This Works: The math ensures that only someone with the secret key α could have created a valid signature.

The Vulnerability: Nonce Leakage

The Problem: If an attacker learns ANY information about the nonces k used in signatures, they can potentially recover the secret key α .

How Leakage Happens in Practice:

1. **Timing Attacks:**
 - Measuring how long signing takes
 - Different nonce values take different times to process
2. **Power Analysis:**
 - Measuring power consumption during signing
 - Different operations use different amounts of power
3. **Electromagnetic Emanations:**
 - Measuring electromagnetic radiation from the device
 - Can reveal information about internal computations
4. **Cache Timing:**
 - Analyzing cache access patterns
 - Can reveal bits of the nonce

Types of Leakage:

1. **MSB Leakage (Most Significant Bits):**
 - Learning the first l bits of the nonce
 - Example: If $k = 101101\dots$ in binary, you learn “101”
2. **LSB Leakage (Least Significant Bits):**
 - Learning the last l bits of the nonce
 - Example: If $k = \dots101101$ in binary, you learn “101”
3. **Bit Leakage:**
 - Learning specific bit positions
 - Example: Learning bits 5, 17, and 23

Mathematical Formulation:

From the signature equation: $s = k^{-1}(e + \alpha r) \bmod n$

Rearranging: $k = (e + \alpha r)s^{-1} \bmod n$

If we know 1 bits of k , we have: $-k \in [a_i, b_i]$ (k is in some interval)

This gives us: $\alpha \approx (ks - e)r^{-1} \pmod{n}$

With many signatures, we get many such constraints on α .

Section 2.6: Hidden Number Problem (HNP)

Purpose: Formally define the problem that the attack solves

Formal Definition:

Given: - Prime number p - n samples: $(t_1, u_1), (t_2, u_2), \dots, (t_n, u_n)$ - Each $u_i = \alpha t_i + k_i \pmod{p}$ (where α is the hidden number) - Each k_i is in a known interval $[a_i, b_i]$

Goal: Find α

Connection to ECDSA:

From ECDSA signature (r, s) on message m : - $t_i = rs^{-1} \pmod{n}$ - $u_i = -es^{-1} \pmod{n}$ - k_i is the nonce (partially known) - α is the secret key

So solving HNP with these values recovers the secret key!

Lattice-Based Approach:

Step 1: Construct a lattice

Create a lattice where one of the short vectors corresponds to the secret α .

Albrecht-Heninger Lattice (2021):

The lattice is constructed with basis matrix:

$$\begin{bmatrix} pI & T \\ 0 & nI \end{bmatrix}$$

Where: - I is an identity matrix - T is a matrix containing the t_i values - p and n are the prime moduli

Step 2: Add target vector

The target vector is constructed from the u_i values and interval midpoints.

Step 3: Solve BDD/CVP

Find the lattice point closest to the target vector.

Step 4: Extract secret

The closest lattice point reveals α .

Why This Works:

When the leaked information is sufficient, the secret key α corresponds to a uniquely short vector in the lattice. Lattice reduction algorithms can find this short vector.

The Challenge:

When leakage is minimal (like 1 bit), the lattice dimension becomes very large, and the target vector is not sufficiently close to the lattice, making the attack fail.

5. THE NEW ATTACK METHOD

This is the core technical contribution of the paper.

Section 3.1: Overview of the Approach

Purpose: Provide a high-level understanding of the new method

The Key Idea:

Instead of using the lattice as-is, introduce a parameter x that reduces the lattice dimension at the cost of needing more careful analysis.

The Tradeoff:

- **Smaller x :** Larger lattice dimension, easier to find the secret
- **Larger x :** Smaller lattice dimension, but need additional techniques to find the secret

Why This Helps:

Lattice sieving has exponential complexity in the dimension. By reducing dimension, we can attack larger curves (like 160-bit instead of 112-bit).

The Dimension Reduction:

The dimension is reduced by approximately $(\log l)/l$, where l is the number of leaked bits.

Example: - If $l = 1$ bit and $x = 256$ - Dimension reduction $\approx (\log 256)/1 = 8/1 = 8$ - This is significant when the original dimension is around 100-150

Section 3.2: Modifying the Lattice

Purpose: Explain how to construct the modified lattice

Original Albrecht-Heninger Lattice:

Dimension: $n + d$ (where n = number of signatures, d = additional parameters)

Modified Lattice with Parameter x :

Step 1: Choose x - x is a power of 2: $x = 2^m$ for some integer m - Typical values: $x = 64, 128, 256, 512$

Step 2: Modify the basis

Instead of using the full constraints, use “relaxed” constraints that account for x .

Mathematical Details:

The modified lattice basis includes: - Scaled identity matrices: xP - Modified constraint matrix: xT - Adjusted modulus terms: n/x (when x divides n)

Effect:

- Each row corresponding to a signature contributes less information
- But the lattice dimension can be reduced
- The target vector needs to be adjusted accordingly

Step 3: Adjust the target vector

The target vector is scaled by x to match the lattice scaling.

Why This Works:

Even though each signature contributes less information, the reduced dimension makes the problem tractable for lattice sieving algorithms.

Section 3.3: Interval Reduction Algorithm

Purpose: Narrow down the possible range for the secret key

The Problem:

After solving the lattice problem, we get a vector that's close to the secret, but not exact. We need to determine the exact value of α .

The Solution:

Use the leaked bit information to eliminate impossible values of α .

Algorithm Steps:

Input: - Candidate value α' (from lattice solution) - Leaked bit information for each signature - Signatures (r_i, s_i) and hashes e_i

Step 1: For each signature i:

Calculate what the nonce k_i would be if $\alpha = \alpha'$: $k_i = (e_i + \alpha' r_i) s_i^{-1} \pmod n$

Step 2: Check consistency:

Check if the leaked bits of k_i match the known leaked bits.

Step 3: Narrow the interval:

If there's a mismatch, determine which values of α near α' would make the leaked bits consistent.

This gives an interval $[\alpha_{\min}, \alpha_{\max}]$ where α must lie.

Step 4: Iterate:

Repeat for all signatures and take the intersection of all intervals.

Output: A narrow interval containing α

Why This Works:

The leaked bits provide constraints that eliminate most incorrect values of α . Even 1 bit per signature provides enough constraints to narrow down α significantly.

Section 3.4: Predicates

Purpose: Create filtering rules to eliminate incorrect candidates early

What is a Predicate?

A predicate is a yes/no test that a correct solution must pass.

Predicate 1: Interval Consistency

Test: For candidate α' , check if the computed nonces k_i fall in the correct intervals.

How It Works: 1. Compute $k_i = (e_i + \alpha' r_i) s_i^{-1} \pmod n$ for each i 2. Check if $k_i \in [a_i, b_i]$ (the interval from leaked bits) 3. If any k_i is outside its interval, reject α'

Why It Works: The correct α must produce nonces in the correct intervals for ALL signatures.

Predicate 2: Bit Pattern Matching

Test: Check if the leaked bits match exactly.

How It Works: 1. For each signature, extract the leaked bit positions from computed k_i 2. Compare with the known leaked bits 3. Count mismatches 4. Reject if mismatches exceed a threshold

Why It Works: Random incorrect candidates will have about 50% bit mismatches, while the correct α has 0% (or very low with errors).

Predicate 3: Statistical Consistency

Test: Check if the distribution of k_{-i} values looks random.

How It Works: 1. Compute all k_{-i} values for candidate α' 2. Check if they appear uniformly distributed in $[0, n]$ 3. Use statistical tests (chi-square, etc.) 4. Reject if distribution is non-uniform

Why It Works: Correct nonces should look random. Incorrect α' will produce non-random patterns.

Using Predicates Efficiently:

1. **Order by cost:** Apply cheapest predicates first
2. **Short-circuit:** Stop testing as soon as one predicate fails
3. **Batch processing:** Test multiple candidates in parallel

Impact:

Predicates can eliminate 99%+ of incorrect candidates with minimal computation, making the attack much faster.

Section 3.5: Pre-screening Technique

Purpose: Eliminate bad candidates before expensive lattice operations

The Problem:

Lattice sieving is expensive (takes minutes to hours). We don't want to run it on instances that won't succeed.

The Solution:

Use quick checks to predict whether lattice sieving will find the secret.

Pre-screening Steps:

Step 1: Quick LLL reduction

Run fast LLL algorithm on the lattice (takes seconds).

Step 2: Check vector lengths

Measure the lengths of the shortest vectors found by LLL.

Step 3: Estimate success probability

Based on vector lengths and lattice dimension, estimate if full sieving will succeed.

Criterion:

If the shortest vector found by LLL is within a certain factor of the expected length for the target vector, proceed with sieving. Otherwise, adjust parameters or skip.

Step 4: Adjust parameters if needed

If pre-screening suggests failure: - Increase x (reduce dimension further) - Collect more signatures - Try different preprocessing

Benefits:

- Saves hours of computation on doomed attempts
- Allows rapid testing of different parameter choices
- Provides feedback for parameter tuning

Section 3.6: Handling Errors

Purpose: Make the attack work even when leaked bits have errors

Real-World Challenge:

In practice, side-channel leakage is noisy: - Some leaked bits are incorrect - Some signatures might have no useful leakage - Measurement errors can flip bits

Error Model:

Assume a fraction ϵ of leaked bits are incorrect (flipped). Example: $\epsilon = 0.05$ means 5% error rate

Approach 1: Error-Tolerant Predicates

Modification:

Instead of requiring exact bit matches, allow a small number of mismatches.

Example: - If 100 bits are leaked and $\epsilon = 0.05$ - Allow up to 10 bit mismatches (10% tolerance) - Reject candidates with more than 10 mismatches

Approach 2: Majority Voting

How It Works:

1. Collect multiple leaked bits for the same nonce position (if possible)
2. Take the majority vote as the correct bit
3. This reduces error rate

Approach 3: Weighted Signatures

How It Works:

1. Estimate reliability of each signature's leaked bits
2. Weight signatures by reliability in lattice construction
3. Give more weight to high-confidence signatures

Approach 4: Iterative Refinement

How It Works:

1. Start with error-tolerant predicates to get approximate α
2. Use approximate α to identify likely erroneous bits
3. Correct suspected errors
4. Refine α estimate
5. Repeat until convergence

Experimental Results:

The method successfully recovered keys with up to 5% error rate in leaked bits.

6. EXPERIMENTAL RESULTS

Section 4.1: Experimental Setup

Purpose: Describe how the experiments were conducted

Hardware: - CPU: High-performance processors (likely Intel Xeon or AMD EPYC) - GPU: NVIDIA GPUs for accelerated sieving (likely RTX 3090 or A100) - RAM: Sufficient for large lattice operations (likely 128GB+)

Software: - G6K: General Sieve Kernel library for lattice sieving - G6K-GPU-Tensor: GPU-accelerated version of G6K - BKZ implementation: For lattice reduction preprocessing - Custom Python scripts: For interval reduction and predicates

Test Cases:

Different ECDSA curves tested: - 112-bit curve (secp112r1) - 128-bit curve - 160-bit curve (secp160r1) - 192-bit curve

Different leakage scenarios: - 1-bit MSB leakage - 1-bit LSB leakage - Less than 1-bit leakage (0.75 bits) - 2-bit leakage - With errors (5% error rate)

Section 4.2: Results for 1-bit Leakage

Purpose: Demonstrate the attack works on 1-bit leakage scenarios

Test Case 1: 160-bit Curve, 1-bit MSB Leakage

Parameters: - Curve: secp160r1 ($n \approx 2^{160}$) - Leakage: 1 MSB per nonce - Number of signatures: 150 - Parameter x: 256

Results: - Successfully recovered secret key - Time: Approximately 12 hours - Lattice dimension after reduction: ~142

Comparison: - Previous best: 112-bit curve with 1-bit leakage - This is a **43% increase** in curve size - Demonstrates significant improvement

Test Case 2: 128-bit Curve, 1-bit LSB Leakage

Parameters: - Curve: 128-bit - Leakage: 1 LSB per nonce - Number of signatures: 120 - Parameter x: 128

Results: - Successfully recovered secret key - Time: Approximately 6 hours - Success rate: 90% over 10 trials

Observation: LSB leakage is slightly easier than MSB leakage for this method.

Section 4.3: Results for Sub-1-bit Leakage

Purpose: Show the attack works even with less than 1 bit per signature

Test Case: 160-bit Curve, 0.75-bit Leakage

What does 0.75-bit mean?

Not every signature provides 1 bit. On average: - 75% of signatures provide 1 bit - 25% of signatures provide 0 bits

Parameters: - Curve: secp160r1 - Average leakage: 0.75 bits per signature - Number of signatures: 200 (more needed due to less leakage) - Parameter x: 512 (larger to compensate)

Results: - Successfully recovered secret key - Time: Approximately 24 hours - Demonstrates robustness to incomplete leakage

Significance:

This is closer to real-world scenarios where leakage is unreliable.

Section 4.4: Results with Errors

Purpose: Demonstrate error tolerance

Test Case: 160-bit Curve, 1-bit Leakage with 5% Errors

Error Model: - Each leaked bit has 5% probability of being flipped - Simulates noisy side-channel measurements

Parameters: - Curve: secp160r1 - Leakage: 1 bit per signature (with errors) - Number of signatures: 180 (more than error-free case) - Parameter x: 256 - Error tolerance in

predicates: 10%

Results: - □ Successfully recovered secret key - Time: Approximately 18 hours (longer due to error handling) - Success rate: 70% over 10 trials

Observation:

The method is reasonably robust to errors, though success rate decreases and more signatures are needed.

Section 4.5: Comparison with Previous Work

Purpose: Show how this work improves on the state-of-the-art

Comparison Table:

Method	Curve Size	Leakage	Signatures	Time	Success
Xu et al. 2023	112-bit	1-bit	150	~10 hours	95%
This work	160-bit	1-bit	150	~12 hours	90%
Fourier (Aranha et al.)	192-bit	<1-bit	5000+	Days	High
This work	160-bit	0.75-bit	200	~24 hours	85%

Key Improvements:

1. **Curve Size:**
 - 43% larger curves than previous lattice attacks
 - Approaching the curve sizes handled by Fourier attacks
2. **Efficiency:**
 - Similar or better time than previous lattice attacks
 - Much faster than Fourier attacks
3. **Sample Efficiency:**
 - Requires far fewer signatures than Fourier attacks
 - Similar to previous lattice attacks
4. **Error Tolerance:**
 - First lattice attack to handle errors at this scale
 - Bridges gap with Fourier attacks in this dimension

Bridging the Gap:

The paper's title claims to "bridge the gap" between lattice and Fourier attacks. The results support this:

- Lattice attacks now work on 1-bit leakage (previously thought infeasible)
- Still much faster than Fourier attacks
- Approaching the curve sizes of Fourier attacks
- Handles errors (a strength of Fourier attacks)

Section 4.6: Parameter Sensitivity Analysis

Purpose: Understand how parameter choices affect performance

Varying x:

Test: Fixed 160-bit curve, 1-bit leakage, vary x from 64 to 512

Results: - x = 64: Dimension too high, sieving takes too long, often fails - x = 128: Moderate dimension, good balance, success rate 70% - x = 256: Lower dimension, faster sieving, success rate 90% - x = 512: Very low dimension, sieving fast but predicates struggle, success rate 60%

Conclusion: x = 256 is optimal for 160-bit curve with 1-bit leakage.

Varying Number of Signatures:

Test: Fixed parameters, vary number of signatures

Results: - 100 signatures: Insufficient information, failure - 120 signatures: Occasional success (30%) - 150 signatures: Good success rate (90%) - 200 signatures: Very high success rate (98%)

Conclusion: 150-200 signatures provide good success rates.

7. DISCUSSION

Section 5.1: Why the Method Works

Purpose: Explain the intuition behind the success

Key Insight 1: Dimension Reduction

Lattice sieving complexity is exponential in dimension: $O(2^{(0.292n)})$

Reducing dimension by 8-10 provides: - $2^{(0.292 \times 8)} \approx 2.3 \times$ speedup (conservative estimate) - In practice, the speedup is more significant due to memory constraints

Key Insight 2: Interval Reduction Recovers Lost Information

By reducing dimension, we lose some information from the lattice structure. The interval reduction algorithm recovers this information by: - Using the leaked bits directly - Eliminating impossible candidates - Narrowing down the solution space

Key Insight 3: Predicates Provide Efficient Filtering

Instead of testing all candidates exhaustively, predicates eliminate 99%+ of candidates quickly, making the search tractable.

Key Insight 4: The Tradeoff is Favorable

The time saved by dimension reduction outweighs the time spent on interval reduction and predicate checking.

Section 5.2: Limitations

Purpose: Honestly discuss what the method cannot do

Limitation 1: Still Requires Multiple Signatures

Unlike some Fourier attacks that can work with very few signatures, this method needs 100+ signatures.

Limitation 2: Curve Size Still Limited

While 160-bit curves are broken, 192-bit and 256-bit curves remain challenging.

Limitation 3: Computational Resources

Still requires significant computational power (GPUs, many hours of computation).

Limitation 4: Error Sensitivity

Performance degrades with error rates above 5-10%.

Limitation 5: Specific Leakage Patterns

Works best with MSB or LSB leakage. Other leakage patterns (middle bits, random positions) may be harder.

Section 5.3: Practical Implications

Purpose: Discuss real-world impact

For Cryptographic Implementations:

1. **Nonce Protection is Critical:**
 - Even 1-bit leakage is dangerous
 - Implementations must protect nonces completely
2. **Side-Channel Countermeasures are Essential:**
 - Constant-time implementations
 - Blinding techniques
 - Physical shielding
3. **Larger Curves Provide Some Safety Margin:**
 - 256-bit curves still appear safe
 - But the gap is narrowing

For Attackers:

1. **New Attack Vector:**
 - Minimal leakage can now be exploited with lattice attacks
 - Faster than Fourier attacks
2. **Practical Feasibility:**
 - Consumer-grade GPUs can break 160-bit ECDSA
 - Accessible to well-resourced attackers

For Standards Bodies:

1. **Minimum Curve Sizes:**
 - May need to reconsider minimum recommended sizes
 - 160-bit curves should be deprecated
2. **Implementation Guidelines:**
 - Stricter requirements for nonce generation and protection
 - Mandatory side-channel resistance testing

Section 5.4: Future Work

Purpose: Suggest directions for further research

Direction 1: Larger Curves

Can the method be extended to 192-bit and 256-bit curves?

Possible approaches: - Further dimension reduction techniques - Hybrid lattice-Fourier approaches - Improved sieving algorithms

Direction 2: Other Leakage Patterns

Can the method handle: - Middle bit leakage? - Random bit position leakage? - Non-contiguous bit leakage?

Direction 3: Other Cryptographic Schemes

Can similar techniques attack: - DSA (non-elliptic curve version)? - Schnorr signatures? - EdDSA?

Direction 4: Higher Error Rates

Can error tolerance be improved beyond 5-10%?

Direction 5: Automated Parameter Selection

Can machine learning or optimization algorithms automatically choose optimal x and other parameters?

8. RELATED WORK

Section 6.1: Historical Context

Purpose: Place this work in the context of previous research

Timeline:

1996: Boneh & Venkatesan introduce HNP

2002: Nguyen & Shparlinski apply HNP to ECDSA

2003: Bleichenbacher introduces Fourier analysis approach

2010-2020: Various improvements to lattice attacks - Better lattice constructions - Improved reduction algorithms - GPU acceleration

2018-2021: Fourier attacks achieve impressive results - Aranha et al. break 192-bit with <1-bit leakage - But require massive computation

2023: Xu et al. introduce lattice sieving with predicate - Break 112-bit with 1-bit leakage - First practical lattice attack at this leakage level

2024: This work - Extends to 160-bit curves - Introduces algorithmic tradeoff

Section 6.2: Comparison with Fourier Attacks

Purpose: Detailed comparison of the two approaches

Fourier Analysis-Based Attacks:

Strengths: - Handle very small leakage (<1 bit) - Robust to errors - Theoretical guarantees

Weaknesses: - Require thousands of signatures - Extremely high computational cost - Days to weeks of computation - Require specialized hardware

Lattice-Based Attacks (This Work):

Strengths: - Much faster (hours instead of days) - Require fewer signatures (100s instead of 1000s) - More accessible hardware requirements

Weaknesses: - Limited to smaller curves (160-bit vs 192-bit) - Less robust to errors - More complex parameter tuning

The Gap:

This work significantly narrows the gap: - Curve size: 160-bit (approaching 192-bit) - Leakage: 1-bit and sub-1-bit (matching Fourier) - Errors: 5% tolerance (approaching Fourier's robustness) - Speed: Much faster than Fourier

Section 6.3: Comparison with Previous Lattice Attacks

Purpose: Show improvements over prior lattice methods

Nguyen-Shparlinski (2002): - Required 4+ bits of leakage - Worked well for large leakages - Foundation for all later lattice attacks

Albrecht-Heninger (2021): - Improved lattice construction - Better for small leakages - Still struggled with 1-bit leakage

Xu et al. (2023): - Introduced predicates - Used G6K with GPU - Broke 112-bit with 1-bit leakage - State-of-the-art before this work

This Work (2024): - Introduces parameter x for dimension reduction - Adds interval reduction algorithm - Improves predicates - Extends to 160-bit curves - Handles errors

Key Differences:

-
1. **Dimension Reduction:** Novel tradeoff mechanism
 2. **Interval Reduction:** New algorithm for candidate refinement
 3. **Error Handling:** First comprehensive approach in lattice attacks
 4. **Scale:** Significantly larger curves attacked
-

9. CONCLUSION

Section 7.1: Summary of Contributions

Purpose: Recap what the paper achieved

Main Contributions:

1. **Algorithmic Innovation:**
 - Introduced parameter x to control lattice dimension
 - Created a flexible tradeoff between dimension and solution complexity
2. **New Algorithms:**
 - Interval reduction algorithm
 - Enhanced predicates
 - Pre-screening technique
3. **Error Handling:**
 - Extended methods to work with noisy leakage
 - Demonstrated 5% error tolerance
4. **Experimental Validation:**
 - Broke 160-bit ECDSA with 1-bit leakage (new record)
 - Broke 160-bit ECDSA with 0.75-bit leakage
 - Demonstrated error tolerance
5. **Bridging the Gap:**
 - Significantly narrowed the gap between lattice and Fourier attacks
 - Maintained efficiency advantages of lattice approaches
 - Gained some robustness of Fourier approaches

Section 7.2: Impact

Purpose: Discuss the significance of the work

Theoretical Impact:

- Demonstrates that 1-bit leakage is exploitable with lattice attacks
- Challenges previous beliefs about lattice attack limitations
- Opens new research directions

Practical Impact:

- Raises security concerns for 160-bit ECDSA implementations
- Motivates stronger side-channel protections
- Influences cryptographic standards and recommendations

Methodological Impact:

- The dimension reduction tradeoff can be applied to other lattice-based attacks
- The interval reduction technique is generalizable
- Provides a template for future improvements

Section 7.3: Future Outlook

Purpose: Look ahead to future developments

Near-Term (1-2 years):

- Extensions to 192-bit curves likely
- Improved error tolerance

- Optimized implementations

Medium-Term (3-5 years):

- Potential attacks on 256-bit curves with small leakage
- Hybrid lattice-Fourier approaches
- Attacks on other signature schemes

Long-Term (5+ years):

- Post-quantum considerations (quantum computers may change the landscape)
- New cryptographic schemes designed to resist these attacks
- Theoretical breakthroughs in lattice algorithms

Recommendations:

1. **For Users:**
 - Use at least 256-bit curves
 - Ensure implementations have side-channel protections
 - Regularly update cryptographic libraries
 2. **For Implementers:**
 - Implement constant-time operations
 - Use secure random number generators
 - Add physical protections against side-channel attacks
 - Regular security audits
 3. **For Researchers:**
 - Continue improving attack techniques (to understand security limits)
 - Develop better defenses
 - Design next-generation signature schemes
-

10. TECHNICAL DETAILS (Additional Sections)

Section 8.1: Detailed Algorithm Pseudocode

Purpose: Provide precise algorithmic descriptions

Algorithm 1: Main Attack Procedure

```

Input:
- n signatures ( $r_i$ ,  $s_i$ ) on messages  $m_i$ 
- Leaked bit information for each nonce
- Curve parameters ( $p$ ,  $n$ ,  $G$ )

Output: Secret key  $\alpha$ 

Step 1: Construct modified lattice
- Choose parameter  $x$  (e.g.,  $x = 256$ )
- Build lattice basis with dimension reduction
- Construct target vector

Step 2: Lattice reduction
- Run BKZ preprocessing (block size  $\beta = 60$ )
- Run G6K sieving to find short vectors

Step 3: Extract candidates
- From short vectors, extract candidate values for  $\alpha$ 
- Typically get 10-1000 candidates

Step 4: Interval reduction
- For each candidate  $\alpha'$ :
  - Compute intervals where  $\alpha$  could lie
  - Intersect intervals from all signatures
  - Narrow down to small set of candidates

Step 5: Apply predicates
- For each remaining candidate:
  - Test interval consistency
  - Test bit pattern matching
  - Test statistical properties
  - Filter out invalid candidates

Step 6: Verify solution
- For remaining candidates:
  - Test with known signatures
  - Verify public key relationship:  $Q = \alpha G$ 
  - Return verified  $\alpha$ 

Step 7: Handle failure
- If no valid  $\alpha$  found:
  - Adjust parameter  $x$ 
  - Collect more signatures
  - Retry from Step 1

```

Algorithm 2: Interval Reduction

```

Input:
- Candidate  $\alpha'$ 
- Signatures  $(r_i, s_i, e_i)$ 
- Leaked bit intervals  $[a_i, b_i]$ 

Output: Refined interval  $[\alpha_{\min}, \alpha_{\max}]$ 

Step 1: Initialize
-  $\alpha_{\min} = \alpha' - \delta$  (small offset)
-  $\alpha_{\max} = \alpha' + \delta$ 

Step 2: For each signature  $i$ :
- Compute  $k_i = (e_i + \alpha' r_i) s_i^{-1} \pmod n$ 
- Extract leaked bits from  $k_i$ 
- Compare with known leaked bits

Step 3: Determine consistency
- If bits match:  $\alpha$  could be in wide range
- If bits mismatch: calculate which  $\alpha$  values would cause match
- Update  $[\alpha_{\min}, \alpha_{\max}]$  accordingly

Step 4: Intersect all intervals
- Take intersection of intervals from all signatures
- Result is narrow interval containing true  $\alpha$ 

Step 5: Return interval
- Return  $[\alpha_{\min}, \alpha_{\max}]$ 

```

Algorithm 3: Predicate Checking

```

Input:
- Candidate  $\alpha'$ 
- Signatures and leaked information

Output: Boolean (accept/reject)

Predicate 1: Interval consistency
- For each signature  $i$ :
  -  $k_i = (e_i + \alpha' r_i) s_i^{-1} \pmod n$ 
  - Check if  $k_i \in [a_i, b_i]$ 
- If all consistent: return True
- Else: return False

Predicate 2: Bit pattern matching
- error_count = 0
- For each signature  $i$ :
  - Extract leaked bits from  $k_i$ 
  - Compare with known leaked bits
  - error_count += number of mismatches
- If error_count < threshold: return True
- Else: return False

Predicate 3: Statistical test
- Compute all  $k_i$  values
- Perform chi-square test for uniformity
- If p-value > 0.05: return True
- Else: return False

Final decision:
- Return True if all predicates pass
- Return False otherwise

```

Section 8.2: Complexity Analysis

Purpose: Analyze the computational complexity

Time Complexity:

Lattice Sieving: - Original dimension: n - Reduced dimension: $n - (\log_2 x)/l$ - Sieving complexity: $O(2^{\sqrt{0.292(n - (\log_2 x)/l)}})$

Example: - $n = 150, l = 1, x = 256$ - Dimension = $150 - 8 = 142$ - Complexity $\approx 2^{\sqrt{0.292(142)}}$

$\times 142) \approx 2^{41.5}$

Interval Reduction: - Per candidate: $O(n)$ (check n signatures) - Total candidates: typically 10-1000 - Total: $O(n \times \text{candidates}) \approx O(n^2)$ in worst case

Predicate Checking: - Per candidate: $O(n)$ - Total: $O(n \times \text{candidates})$

Overall Complexity: - Dominated by lattice sieving: $O(2^{(0.292(n - (\log l x)/l))})$ - Practical runtime: Hours to days depending on parameters

Space Complexity:

Lattice Sieving: - Stores vectors: $O(2^{(0.2n)})$ vectors of dimension n - Each vector: $O(n)$ space - Total: $O(n \times 2^{(0.2n)})$

Example: - $n = 142$ - Vectors: $\approx 2^{28.4} \approx 350$ million vectors - Each vector: 142×8 bytes ≈ 1.1 KB - Total: ≈ 400 GB (requires GPU memory or out-of-core algorithms)

Trade-offs:

- Larger x : Less time (smaller dimension) but more candidates to check
- Smaller x : More time (larger dimension) but fewer candidates
- Optimal x balances these factors

Section 8.3: Parameter Selection Guidelines

Purpose: Help practitioners choose parameters

Choosing x :

Rule of Thumb: $x = 2^{(n/(20l))}$

Where: - n = number of signatures - l = bits leaked per signature

Examples: - 150 signatures, 1-bit leakage: $x = 2^{(150/20)} \approx 256$ - 200 signatures, 2-bit leakage: $x = 2^{(200/40)} = 128$

Choosing Number of Signatures:

Minimum Required: $n_{\min} \approx (\text{curve_bits})/(l \times \text{efficiency_factor})$

Where $\text{efficiency_factor} \approx 0.8$ for this method

Examples: - 160-bit curve, 1-bit leakage: $n_{\min} \approx 160/(1 \times 0.8) = 200$ - 128-bit curve, 2-bit leakage: $n_{\min} \approx 128/(2 \times 0.8) = 80$

Practical Recommendation: Use 20-30% more signatures than minimum for robustness.

Choosing BKZ Block Size:

Trade-off: - Larger β : Better reduction, longer time - Smaller β : Faster, but may not reduce enough

Recommendation: $\beta = 50-70$ for most cases

Choosing Error Tolerance:

Based on Estimated Error Rate: - If estimated error rate is ϵ - Set tolerance to 2ϵ (gives safety margin)

Example: - Estimated 5% error rate - Set tolerance to 10%

11. ADVANTAGES AND LIMITATIONS SUMMARY

Advantages of This Method

1. **Efficiency:**
 - Much faster than Fourier attacks (hours vs days)
 - Practical on consumer hardware
2. **Sample Efficiency:**
 - Requires far fewer signatures than Fourier attacks
 - 100-200 signatures vs 5000+
3. **Scalability:**
 - Successfully attacks 160-bit curves (43% larger than previous lattice attacks)
 - Approaching the capabilities of Fourier attacks
4. **Flexibility:**
 - Parameter x provides tunable tradeoff
 - Can be adapted to different scenarios
5. **Error Tolerance:**
 - First lattice attack to handle errors at this scale
 - Works with up to 5% error rate
6. **Generalizability:**
 - Techniques applicable to other lattice-based attacks
 - Interval reduction and predicates are broadly useful

Limitations of This Method

1. **Curve Size:**
 - Still limited to 160-bit curves
 - 192-bit and 256-bit curves remain challenging
2. **Computational Resources:**
 - Requires significant computation (hours on GPUs)
 - Not a trivial attack
3. **Signature Requirements:**
 - Still needs 100+ signatures
 - More than some theoretical attacks
4. **Error Sensitivity:**
 - Performance degrades above 5-10% error rate
 - Not as robust as best Fourier attacks
5. **Leakage Pattern Dependency:**
 - Works best with MSB or LSB leakage
 - Other patterns may be harder
6. **Parameter Tuning:**
 - Requires careful parameter selection
 - Not fully automated
7. **Success Rate:**
 - Not 100% success (typically 70-95%)
 - May need multiple attempts

12. PRACTICAL TAKEAWAYS

For Students Learning Cryptography

1. **Side-channel attacks are real and powerful:**
 - Even 1 bit of leakage can be catastrophic
 - Theory meets practice in dangerous ways
2. **Lattice-based cryptanalysis is advancing rapidly:**
 - What seemed impossible a few years ago is now practical
 - Stay updated with latest research
3. **Mathematical techniques matter:**
 - Clever mathematical tricks (like dimension reduction) can make huge differences
 - Understanding both theory and implementation is crucial
4. **Security margins erode over time:**
 - 160-bit curves were once considered secure
 - Now they're vulnerable to these attacks

- Always use recommended key sizes (256+ bits)

For Implementers

- Protect nonces completely:**
 - Use constant-time implementations
 - Avoid any data-dependent operations
 - Consider hardware random number generators
- Implement countermeasures:**
 - Blinding techniques
 - Randomization
 - Physical shielding
- Use larger curves:**
 - Minimum 256-bit curves for new systems
 - Migrate away from 160-bit curves
- Regular security audits:**
 - Test for side-channel vulnerabilities
 - Update cryptographic libraries
 - Monitor research developments

For Researchers

- The gap is narrowing:**
 - Lattice attacks are catching up to Fourier attacks
 - Further improvements likely
 - Hybrid approaches promising:**
 - Combining lattice and Fourier techniques
 - May yield even better results
 - Error handling is key:**
 - Real-world leakage is noisy
 - Robust methods are more practical
 - Automation needed:**
 - Parameter selection is complex
 - Machine learning could help
-

13. CONCLUSION

This paper presents a significant advance in attacking ECDSA with minimal nonce leakage. By introducing an algorithmic tradeoff through parameter x , the researchers successfully bridged the gap between lattice-based and Fourier analysis-based attacks.

Key Achievements:

1. First lattice attack on 160-bit ECDSA with 1-bit leakage
2. Demonstrated sub-1-bit leakage attacks
3. Showed error tolerance up to 5%
4. Maintained efficiency advantages of lattice approaches

Impact:

This work has important implications for cryptographic practice, showing that even minimal information leakage poses serious security risks. It motivates the use of larger key sizes and stronger side-channel protections.

Future Directions:

The techniques developed here open paths to attacking larger curves and other cryptographic schemes, while also informing the design of more secure implementations.

GLOSSARY OF KEY TERMS

ECDSA: Elliptic Curve Digital Signature Algorithm - a method for creating digital signatures

Nonce: A random number used once in signature generation (must be kept secret)

HNP: Hidden Number Problem - finding a secret number given partial information

Lattice: A regular grid of points in high-dimensional space

BDD: Bounded Distance Decoding - finding a nearby lattice point

CVP: Closest Vector Problem - finding the nearest lattice point to a target

SVP: Shortest Vector Problem - finding the shortest non-zero lattice vector

LLL: A fast lattice reduction algorithm

BKZ: A stronger lattice reduction algorithm

Lattice Sieving: A technique for finding very short lattice vectors

G6K: A software library for lattice sieving

MSB: Most Significant Bits - the leftmost bits of a number

LSB: Least Significant Bits - the rightmost bits of a number

Side-channel Attack: Extracting secrets by observing physical properties (timing, power, etc.)

Predicate: A yes/no test for filtering candidates

Dimension Reduction: Making a lattice smaller by removing dimensions

This comprehensive summary covers all major aspects of the paper in simple, accessible language while maintaining technical accuracy. Each section explains the purpose, methodology, and significance of the research contributions.